






[PO]MDP in practice








# The tiger problem




# Outcomes

		State	
			
Action		-100	+ 10
		+ 10	-100
		-1	-1

# Observations

			
		0.50	0.50
		0.50	0.50
		<b>0.85</b>	0.15
		0.50	0.50
		0.50	0.50
		0.15	<b>0.85</b>

# Transitions



# POMDP toolbox

- Download from  
Official [www.pomp.org](http://www.pomp.org) <http://cs.brown.edu/research/ai/pomdp/>  
Github <https://github.com/lionel-rigoux/CPC-pomdp>
- Compile the C code
- Edit the POMDP description file
- Run the executable or the the helper script




# POMDP solution

## alpha vectors (.alpha file)

```
// action  
// v1 v2 v3 ...  
//  
// ...
```

$b = [p(s=s_1) \ p(s=s_2) \ p(s=s_3) \ \dots]$ , sum to 1

$V(\text{action}, b) = v_1 \ p(s=s_1) + v_2 \ p(s=s_2) + v_3 \ p(s=s_3)$



Some actions can have multiple alpha vectors, some none!



# POMDP simulation

## internalized world

$T(s' | s, a)$ ,  $O(o | s, a)$ ,  $R(a, s)$



## real world

$T(s' | s, a)$ ,  $O(o | s, a)$ ,  $R(a, s)$



*INITIALIZE*

- a. find optimal policy
- b. pick initial belief

- a. pick initial state

*ITERATE*

1. chose best action given current belief

a

2. change state depending on action
3. get reward depending on state and action
4. get observation depending on state and action

o

5. update belief following given action, observation and internal model



# POMDP summary

## **The toolbox allows to:**

- Compute the optimal policy for a given representation of the world
- Compute the utility/value associated with beliefs and actions
- Simulate the effect of a false representation of the world on complex decision making and belief updating

## **It can not:**

- capture how we learn the policy
- capture how representation of the world are learned
- be used to fit experimental data (infer on hidden representations)




# Questions?






# POMDP solution

belief MDP (.pg file)



$$p(o_{TL} \mid \text{listen}, TL) = 0.99$$




$$p(o_{TL} \mid \text{listen}, TL) = 0.75$$




# Resolution

immediate reward



future value  $V(b') = \sum V(b'|o)p(o|b)$



belief value  $V(b) = r(b) + \gamma V(b')$

